

# Stroke Prediction

Group 10

Jared Clarke & Amr Salem

A dark blue diagonal gradient bar that starts from the bottom left corner and extends towards the top right corner, covering the lower half of the slide.

# Why strokes?


- Strokes are the 2nd leading cause of death in developed countries after heart disease (WHO)
- The #5 cause of death in the US
- One of the leading causes of disability
- Sometimes can be prevented
- Could be of use medically




**American  
Stroke  
Association.**

*A division of the  
American Heart Association.*


# What is a stroke?



- Blood vessel that supplies nutrients to the brain is hindered by a clot (Ischemic -- 87 % of strokes)
- Blood vessel bursts (Hemorrhagic)
- Brain tissues can't receive necessary nutrients
- Leads to brain cells dying in only a few minutes
- Difficulty walking, speaking and understanding, paralysis of the face and or extremities
- Deadly if not treated immediately.



# Stroke Risk Factors



- Age -- Especially for people > 65
- Gender -- women > men
- Hypertension
- High Cholesterol
- Smoking
- Diabetes
- Obesity - high BMI (over 30)
- Stress (work, marriage, location of residence)

# Data Sources:

Kaggle



Analytics Vidhya



McKinsey Analytics  
Online Hackathon  
Healthcare Analytics



McKinsey Analytics Online Hackathon - Healthcare Analytics

Question we hope  
to answer:

Can we reliably predict  
a stroke based on  
certain features of a  
person's medical  
history?

# Features

- Gender (Male, Female, Unknown)
- Age
- Hypertension (high blood pressure)
- Heart Disease (yes or no)
- Has ever been married (yes or no)
- Work type (Private, self-employed, government job, etc.)
- Residence type (Urban or Rural)
- Avg glucose level
- BMI (>30 considered high risk)
- Smoking status
- Has patient ever experienced a stroke

# Processing the data

```
# Import data from PostgreSQL database
db_string = f"postgres://postgres:{db_password}@127.0.0.1:5432/strokes_db"
engine = create_engine(db_string)
stroke_df = pd.read_sql_table('total_stroke_data', con=engine)

stroke_df.head()
```

```
:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	18069	Male	70.0	1	0	Yes	Self-employed	Urban	104.24	34.7	formerly smoked	0
1	49086	Female	23.0	0	0	No	Private	Urban	60.50	27.1	formerly smoked	0
2	19671	Female	58.0	0	0	Yes	Govt_job	Urban	93.15	34.7	never smoked	0
3	59225	Male	48.0	1	0	Yes	Govt_job	Urban	55.25	49.7	never smoked	0
4	25175	Female	56.0	0	0	No	Private	Rural	108.50	28.0	never smoked	0

```
stroke_df.shape
```

```
(30555, 12)
```

```
stroke_df['ever_married'] = stroke_df['ever_married'].apply(lambda x: 1 if x == 'Yes' else 0)
stroke_df['residence_type'] = stroke_df['residence_type'].apply(lambda x: 1 if x == 'Urban' else 0)
```

```
# Encoding the gender column
```

```
gender_num = []
for i in stroke_df['gender']:
    if i == 'Male':
        gender_num.append(0)
    if i == 'Female':
        gender_num.append(1)
```

```
stroke_df['gender'] = gender_num
```

```
# Encoding the 'work_type' column
```

```
label_encoder = LabelEncoder()
label_encoder.fit(stroke_df['work_type'])
stroke_df['work_type_le'] = label_encoder.transform(stroke_df['work_type'])
```

```
work_type_num = {'Private': 0,
                  'Self-employed': 1,
                  'Govt_job': 2,
                  'children': 3,
                  'Never_worked': 4}
```


```
stroke_df['work_type_num'] = stroke_df['work_type'].apply(lambda x: work_type_num[x])
stroke_df.drop(columns=['work_type', 'work_type_le'], inplace=True)
```



# Exploring the data:




- Nearly 2% of people in the dataset had a stroke
- Slight gender imbalance toward women
- Average age is 68.35 years

## Total Stroke Percentage




- Higher number of strokes for former smokers compared to current smokers (180 : 112)
- 28% or 162 people had hypertension
- 22% or 125 people had heart disease
- 329 (57%) Private company employers and 176 self-employed (31%) had strokes



Stroke/Smoking Data



Age vs avg\_glucose\_level




bmi vs avg\_glucose\_level



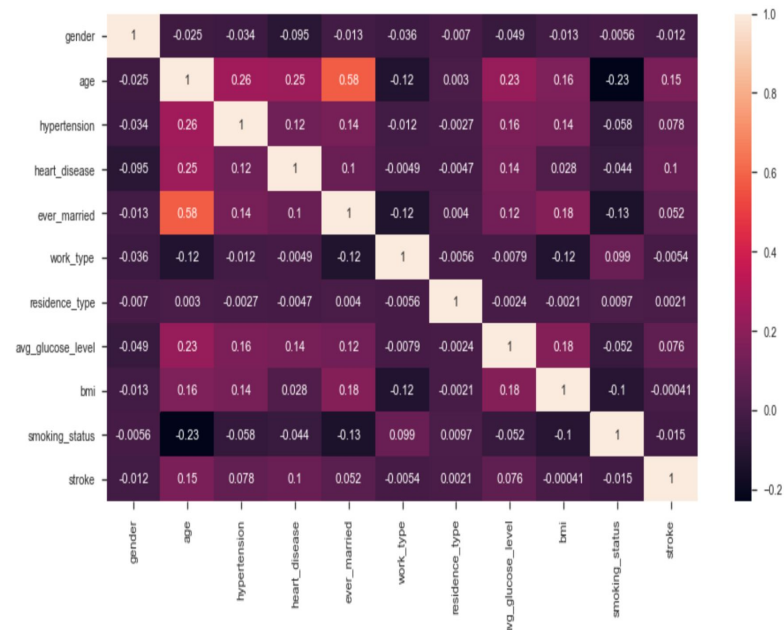
Ever Married

No

Yes



# Correlation



Residuals:

	Min	1Q	Median	3Q	Max
	-0.13081	-0.03178	-0.01290	0.00119	1.01200

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-2.610e-02	4.376e-03	-5.964	2.48e-09 ***
gender	9.339e-06	1.580e-03	0.006	0.9953
age	1.094e-03	5.217e-05	20.973	< 2e-16 ***
hypertension	1.543e-02	2.588e-03	5.961	2.53e-09 ***
heart_disease	3.911e-02	3.636e-03	10.756	< 2e-16 ***
ever_married	-1.424e-02	2.136e-03	-6.669	2.61e-11 ***
work_type	4.570e-04	8.680e-04	0.526	0.5986
residence_type	5.570e-04	1.532e-03	0.364	0.7161
avg_glucose_level	1.142e-04	1.789e-05	6.383	1.76e-10 ***
bmi	-5.194e-04	1.093e-04	-4.752	2.02e-06 ***
smoking_status	3.156e-03	1.000e-03	3.156	0.0016 **

---  
 Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1339 on 30537 degrees of freedom  
 Multiple R-squared: 0.03353, Adjusted R-squared: 0.03322  
 F-statistic: 106 on 10 and 30537 DF, p-value: < 2.2e-16

# Analysis

- Random Forest Classifier
- Logistic Regression

## Random Forest Classifier

```
forest = RandomForestClassifier(n_estimators = 100)
forest.fit(X_train, y_train)

forest_score = forest.score(X_train, y_train)
forest_test = forest.score(X_test, y_test)

y_pred = forest.predict(X_test)

print('Training Score', forest_score)
print('Testing Score \n', forest_test)
print(cm)
print(classification_report(y_test, y_pred))
```

Training Score 0.9993777777777778  
Testing Score  
0.8173333333333334  
[[5848 0]  
 [1652 0]]

		precision	recall	f1-score	support
	0	0.84	0.94	0.89	5848
	1	0.65	0.37	0.47	1652
	accuracy			0.82	7500
	macro avg	0.75	0.66	0.68	7500
	weighted avg	0.80	0.82	0.80	7500

## ## Logistic Regression

```
model = LogisticRegression(solver="lbfgs", max_iter=200)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

print('Testing Score \n', score)

print(classification_report(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)
print(cm)
```

Testing Score  
0.7797333333333333

		precision	recall	f1-score	support
	0	0.78	1.00	0.88	5848
	1	0.00	0.00	0.00	1652
	accuracy			0.78	7500
	macro avg	0.39	0.50	0.44	7500
	weighted avg	0.61	0.78	0.68	7500

[[5848 0]  
 [1652 0]]

# Neural Network

239/239 - 1s - loss: 0.0930 - accuracy: 0.9811  
Loss: 0.09303482621908188, Accuracy: 0.9811444282531738

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 80)	880
dense_1 (Dense)	(None, 30)	2430
dense_2 (Dense)	(None, 1)	31

Total params: 3,341

Trainable params: 3,341

Non-trainable params: 0

# Stroke Risks You Can Control:

- High Blood Pressure
- Smoking
- Diabetes
- Diet
- Physical Inactivity
- Cholesterol
- Obesity

## Unique Symptoms in Women:

- Loss of consciousness or fainting
- General weakness
- Difficulty or shortness of breath
- Confusion, unresponsiveness or disorientation
- Sudden behavioral change
- Agitation
- Hallucination
- Nausea or vomiting
- Pain
- Seizures
- Hiccups

# Takeaways

- Too much data for Logistic Regression
- Unbalance data set -- 1.9% had strokes
- Leave out features that are weak according to the heatmap and test
- Leave out features that are weak according to multiple linear regression and test
- Want precision more than accuracy



# Analysis Tools Used

- GitHub
- Python
- Jupyter Notebook
- PostgreSQL
- Scikit learn library
- Keras library
- Tensorflow
- Tableau
- SQLAlchemy
- Visual Studio Code